

# The Tolman Mechanism

Development Notes

Stefano Ghirlanda

December 7, 2020

## 1 Rationale

- The Tolman mechanism implements a learning mechanism that formalizes common theorizing in animal cognition.
- The animal is assumed to build a mental model of its environment.
- The model is then used to plan how to reach states with high value.
- There are major computational challenges in such a mechanism and it is unlikely that animals (or even people) can apply this strategy in cases beyond a certain complexity.
- We restrict the mechanism to work in worlds with a single goal, and that “restart” after the goal is obtained.

## 2 Variables

- The Tolman mechanism learns  $S \rightarrow B \rightarrow S'$  transition probabilities.
- These estimates are called  $z(S, B, S')$ .
- Everything else is decided when responding.

## 3 Learning

- There is a single learning rate,  $\alpha_z$ .

- Imagine first that all stimuli are made of a single element. When  $S \rightarrow B \rightarrow S'$  is observed,  $z(S, B, X)$  is updated for all  $X$  as follows:

$$\Delta z(S, B, X) = \alpha_z (\lambda_X - z(S, B, X)) \quad (1)$$

where  $\lambda_X = 1$  when  $X = S'$  (the state that actually occurred), and 0 for all other states.

- With sufficient experience and sufficiently small  $\alpha_z$ , eq. (3) will converge to the actual transition probabilities.
- How to extend to stimuli with more elements? Let's first consider that all intensities are = 1. We can use

$$\forall S'_j : \quad z(S, B, S'_j) = \sum_{i=1}^n z(S_i, B, S'_j) \quad (2)$$

where the  $S_i$  are the elements of  $S$  and the  $S'_j$  the elements of  $S'$ . According to eq. (2), different elements of  $S$  compete for predicting each element of  $S'$  in the same way as they would compete for accruing  $v$  or  $w$  values in A-learning (and other mechanisms).

- Using eq. (2) to calculate  $z(S_i \dots S_n, B, S'_i)$  we can update each  $S_i$  as follows:

$$\Delta z(S_i, B, X) = \alpha_z (\lambda_X - z(S, B, S'_j)) \quad (3)$$

where  $\lambda = 1$  if  $X \in S'$  and  $\lambda = 0$  otherwise.

- With intensities written as  $x$  we can use:

$$\forall S'_j : \quad z(S, B, S'_j) = \sum_{i=1}^n z(S_i, B, S'_j) x_i \quad (4)$$

$$\Delta z(S_i, B, S'_j) = \alpha_z (\lambda_X - z(S, B, S'_j)) x_i x'_j \quad (5)$$

### 3.1 Trial implementation <2020-12-07 Mon>

- The equations above have been implemented in the `github` branch `tolman-mechanism`
- I have made a super simple test where there is only behavior  $B$  and two stimuli  $S_1$  and  $S_2$ , and the only thing that happens is the sequence:

$$\dots \rightarrow S_1 \rightarrow B \rightarrow S_2 \rightarrow B \rightarrow S_1 \rightarrow \dots \quad (6)$$

- In this case, the Tolman mechanism should learn:


$$z(S_1, B, S_2) = 1 \quad (7)$$

$$z(S_1, B, S_1) = 0 \quad (8)$$

$$z(S_2, B, S_1) = 1 \quad (9)$$

$$z(S_2, B, S_2) = 0 \quad (10)$$

- Printing this values throughout the simulation shows that they are in fact learned, but the graphs do not come out well. I think this is because the current plotting code only updates the plot for stimulus elements that are actually presented, but the Tolman mechanism also updates  $z(S, B, S')$  for  $S'$  elements that are not present. For example, this is a graph I get:



- The initial values of all  $z$  is 0.1 for testing purposes. The blue line goes to 1 as it should. The orange is plotted at the initial 0.1 throughout the simulation, but for the last point, which is the correct value of 0. As I mentioned above, however, if I print the values during learning I do see that they are decreasing toward 0 all the time.
- (Element intensities and multiple elements are implemented but not tested yet.)

## 4 Responding

- Responding assumes that the world model in  $z$  is true.

- Responding should then make the “best plan” given this knowledge.
- Evaluating all possible plans is tricky in general. Let’s work with a particular kind of world to begin with:
  1. There is only one valued stimulus, called the goal.
  2. When the goal is reached, the next state is determined by the world using a fixed rule (deterministic or stochastic).
- The second condition means that we need to consider just how to reach the goal once, since after that the world “resets” to a statistically equivalent state.
- The first step is to find all paths to the goal. We start from the goal and we go backwards along possible transitions, that is, all transitions with  $z > 0$  for some  $B$ . If we eventually reach the current state, we store the sequence of transitions and call it a “path.”
- We calculate the expected *rate* of return for all paths to the goal, that is the expected value divided the path length (number of actions):

$$v(\text{path}) = \frac{u(S_{\text{goal}}) - \sum_{B \in \text{path}} c(B)}{l_{\text{path}}} \prod_{S' \in \text{path}} \Pr(S \rightarrow B \rightarrow S') \quad (11)$$

- In eq. (11), the sum is over all behaviors in the path,  $l_{\text{path}}$  is path length, and the product is over all transition probabilities in the path. (This is the probability that the plan will succeed.)
- We then give a value to each behavior that is feasible in response to the current stimulus. If the behavior is the first step on a path to the goal, its value is the  $v$  value of the goal. If it is not, its value is `start_v`.
- If a behavior is the first on more than one path, we can average the  $v$ ’s using the success probability of each path, so we should store these somewhere.
- We then use softmax to choose a behavior.