

The Tolman Mechanism

Development Notes

Stefano Ghirlanda

December 7, 2020

1 Rationale

- The Tolman mechanism implements a learning mechanism that formalizes common theorizing in animal cognition.
- The animal is assumed to build a mental model of its environment.
- The model is then used to plan how to reach states with high value.
- There are major computational challenges in such a mechanism and it is unlikely that animals (or even people) can apply this strategy in cases beyond a certain complexity.
- We restrict the mechanism to work in worlds with a single goal, and that “restart” after the goal is obtained.

2 Variables

- The Tolman mechanism learns $S \rightarrow B \rightarrow S'$ transition probabilities.
- These estimates are called $z(S, B, S')$.
- Everything else is decided when responding.

3 Learning

- There is a single learning rate, α_z .

- Imagine first that all stimuli are made of a single element. When $S \rightarrow B \rightarrow S'$ is observed, $z(S, B, X)$ is updated for all X as follows:

$$\Delta z(S, B, X) = \alpha_z (\lambda_X - z(S, B, X)) \quad (1)$$

where $\lambda_X = 1$ when $X = S'$ (the state that actually occurred), and 0 for all other states.

- With sufficient experience and sufficiently small α_z , eq. (3) will converge to the actual transition probabilities.
- How to extend to stimuli with more elements? Let's first consider that all intensities are = 1. We can use

$$\forall S'_j : \quad z(S, B, S'_j) = \sum_{i=1}^n z(S_i, B, S'_j) \quad (2)$$

where the S_i are the elements of S and the S'_j the elements of S' . According to eq. (2), different elements of S compete for predicting each element of S' in the same way as they would compete for accruing v or w values in A-learning (and other mechanisms).

- Using eq. (2) to calculate $z(S_i \dots S_n, B, S'_i)$ we can update each S_i as follows:

$$\Delta z(S_i, B, X) = \alpha_z (\lambda_X - z(S, B, S'_j)) \quad (3)$$

where $\lambda = 1$ if $X \in S'$ and $\lambda = 0$ otherwise.

- With intensities written as x we can use:

$$\forall S'_j : \quad z(S, B, S'_j) = \sum_{i=1}^n z(S_i, B, S'_j) x_i \quad (4)$$

$$\Delta z(S_i, B, S'_j) = \alpha_z (\lambda_X - z(S, B, S'_j)) x_i x'_j \quad (5)$$

4 Responding

- Responding assumes that the world model in z is true.
- Responding should then make the “best plan” given this knowledge.
- Evaluating all possible plans is tricky in general. Let's work with a particular kind of world to begin with:

1. There is only one valued stimulus, called the goal.
 2. When the goal is reached, the next state is determined by the world using a fixed rule (deterministic or stochastic).
- The second condition means that we need to consider just how to reach the goal once, since after that the world “resets” to a statistically equivalent state.
 - The first step is to find all paths to the goal. We start from the goal and we go backwards along possible transitions, that is, all transitions with $z > 0$ for some B . If we eventually reach the current state, we store the sequence of transitions and call it a “path.”
 - We calculate the expected *rate* of return for all paths to the goal, that is the expected value divided the path length (number of actions):

$$v(\text{path}) = \frac{u(S_{\text{goal}}) - \sum_{B \in \text{path}} c(B)}{l_{\text{path}}} \prod_{S' \in \text{path}} \Pr(S \rightarrow B \rightarrow S') \quad (6)$$

- In eq. (6), the sum is over all behaviors in the path, l_{path} is path length, and the product is over all transition probabilities in the path. (This is the probability that the plan will succeed.)
- We then give a value to each behavior that is feasible in response to the current stimulus. If the behavior is the first step on a path to the goal, its value is the v value of the goal. If it is not, its value is **start_v**.
- If a behavior is the first on more than one path, we can average the v ’s using the success probability of each path, so we should store these somewhere.
- We then use softmax to choose a behavior.